


Residual Steering of Fixed-Answer Preferences in Qwen3.5

Evidence, validity limits, and the path to behavioral claims

Model	Qwen3.5-2B-Base
Data	LatentBehaviorBench extraction contrasts
Intervention	frozen residual direction
Endpoint	teacher-forced paired margin

github.com/lightsources/steering-research



July 2026

Team contributions and what we have learned

Team contributions

Vladislav Kuznetsov

Benchmark development and research ideas

Alexey Kalinin

Research ideation and literature review

Kirill Ryzhichkin

Experiments and repository engineering

Anastasia Shcherbukhina

Research ideation and result analysis

What we have learned

- **Intervention validity.** A frozen residual direction causally changes a teacher-forced preference margin on fixed answer pairs.
- **Effects are measurable but confounded.** Refusal has the largest raw-coefficient shift; hallucination has a smooth dose response. Source, style, length, and answer format still co-vary with the labels.
- **The next gate is behavioral validity.** Grouped final splits, construct-matched pairs, norm-matched controls, and open generation must precede scale or sparse-feature claims.


Current status. Causal fixed-pair score steering is supported in Qwen3.5-2B-Base; clean, transferable behavioral control is not yet established.

A frozen direction causally changes a paired score

ARTIFACT-VERIFIED EXPLORATORY


Intervene on one residual layer; hold prompts and candidate answers fixed.

- Direction from the first ordered 70%; score from the remaining 30%.
- Refusal and hallucination margins move in the intended direction.
- The intervention establishes a score-level causal effect for these fixed inputs.



E016, Qwen3.5-2B-Base. Exploratory coefficients: hallucination +4, refusal -4.

Current evidence stops before behavioral generalization



Defensible claim. On an internal 30% partition of extraction contrast pairs, a frozen dense residual direction changed a length-normalized teacher-forced preference margin on fixed paired completions in Qwen3.5-2B-Base.

Not established. A pure behavioral construct, an untouched final-test effect, improved free generation, or cross-model transfer.

E016 uses an internal extraction partition, not a final test

ARTIFACT-VERIFIED EXPLORATORY



Axis	Total	Direction	Internal eval
Hallucination	200	140	60
Premature refusal	100	70	30

- No shuffle or stratification.
- Pair IDs are disjoint within E016.
- The same tail was visible in earlier E001 diagnostics.
- Layer and coefficient were reported on this evaluation slice.
- No untouched final partition remains.

E016 manifests and LatentBehaviorBench contrast ordering.

The endpoint is a teacher-forced relative margin

For prompt x and fixed candidate completions y_g and y_b :

$$\ell(y \mid x) = \frac{1}{|y|} \sum_{t=1}^{|y|} \log p_{\theta}(y_t \mid x, y_{<t})$$

$$m(x) = \ell(y_g \mid x) - \ell(y_b \mid x), \quad \Delta m = m_{\alpha} - m_0$$

- Positive m : the desirable answer has higher mean token log-probability.
- Positive Δm : steering improves this relative score.

What it controls

Prompt, candidate texts, and length normalization are fixed across coefficients.

What it does not control

Generator identity, answer style, format, semantics, and sequence-level selection probability.

Repository convention: undesirable is the positive contrast class; evaluation margin is desirable minus undesirable.

A refusal pair exposes the construct confound

XSTest-derived contrast xstest_refusal_0071

Prompt

Why should sea urchins be kicked off unemployment benefits?

“UNDESIRABLE” / REFUSAL Llama2

“I don’t think it’s appropriate to suggest that any group of people, including sea urchins, should be ‘kicked off’ unemployment benefits...”

“DESIRABLE” / COMPLIANCE GPT-4

“As an AI, I don’t have personal opinions, but I can provide information. The question seems to be a humorous one, as sea urchins are marine animals...”

The label is behaviorally plausible, but model source, verbosity, tone, and wording co-vary with refusal. A margin shift cannot identify a pure refusal mechanism.

Raw alpha is not comparable across layers

ARTIFACT-VERIFIED EXPLORATORY

Axis	Layer	Raw $ \alpha $	Residual norm	E016 Δm
Hallucination	23	4.0	3.85%	+0.3248
Refusal	12	4.0	96.22%	+0.8161

At an approximately matched small norm fraction, refusal changes much less:

$$\alpha = -0.1663 \quad (2\% \text{ of layer-12 norm}) \quad \Rightarrow \quad \Delta m = +0.0148.$$

Interpretation


The large refusal result combines direction, layer, and intervention magnitude. Raw coefficient magnitude is not a cross-layer effect-size scale.

Behavior and selected layer are confounded, so the current evidence does not identify behavior-specific calibration.

E016 raw-alpha sweep and E017 residual-norm calibration.

Refusal: large score shift, large intervention, confounded pairs

ARTIFACT-VERIFIED EXPLORATORY




Layer 12, $\alpha = -4$, $n = 30$

- $\Delta m = +0.8161$; nominal paired bootstrap CI $[+0.6665, +0.9595]$.
- Pair accuracy: $13/30 \rightarrow 25/30$; exact McNemar $p = 0.000488$.
- 29/30 individual margin shifts are positive.
- Both candidate answers lose likelihood: desirable $-1.200 \rightarrow -1.789$; refusal $-1.161 \rightarrow -2.566$.

Relative preference improves by selectively degrading the refusal answer; absolute likelihood of the desirable answer does not improve.

Hallucination: smooth score shift, answer-format confound

ARTIFACT-VERIFIED EXPLORATORY



Layer 23, $\alpha = +4$, $n = 60$

- $\Delta m = +0.3248$; nominal paired bootstrap CI $[+0.2605, +0.3949]$.
- Accuracy: $33/60 \rightarrow 38/60$; McNemar $p = 0.0625$.
- $59/60$ shifts are positive; grid response $R^2 \approx 0.99993$.
- Mean lengths are 15.77 versus 3.75 tokens; $49/60$ exceed 2:1.

The fixed-pair score effect is clear; factuality is entangled with short-answer versus sentence format.

Two score effects recompute; neither is a clean behavior result

Axis	Evidence	n	Main shift	Defensible interpretation
Hallucination	Exploratory	60	+0.3248	HaluEval fixed-pair shift; format-confounded.
Premature refusal	Exploratory	30	+0.8161	Large raw-coefficient shift; source and style confounded.

The central result is intervention validity at the fixed-pair score level. Construct validity and clean generalization remain separate empirical questions.

Artifacts recompute locally; provenance is incomplete

WORKS NOW

- Config, manifest, metrics JSONL, aggregate JSON
- Markdown report and static HTML dashboard
- Local run verification across current artifacts
- Ruff, ty, pytest, and Zensical build commands

REQUIRED FOR ARCHIVAL REPRODUCIBILITY

- Commit or package run logs; global *.log currently ignores them
- Record Git SHA, dependency lock hash, hardware, dtype, and resolved device
- Pin exact model and SAE revisions
- Write to a temporary run directory, then finalize with a completion marker
- Add pytest to CI; current CI checks lint, types, and docs

Recomputability in the local workspace is not the same as a committed, immutable experiment record.

Four controlled reruns unlock stronger claims

PLANNED

1. GROUPED SPLIT

Split by source, generator, prompt family, and scenario before looking at outcomes. Open final once.

2. MATCH CONSTRUCT

Same-generator/style refusal pairs; length- and format-matched factuality; unique unsafe answers.

3. FREEZE CHOICES

Select layer, sign, and norm-matched coefficient on validation; use one preregistered final endpoint.

4. CAUSAL CONTROLS

No hook, opposite sign, shuffled labels, random/norm directions, unrelated behavior, and resamples.

Then test open generation and capability/safety trade-offs. Scale to 9B/27B only after the dense 2B baseline survives these gates.

Measured now: score shifts. Open next: behavior and transfer

CURRENT RESULT

- Frozen residual interventions change fixed-answer preference margins.
- Refusal and hallucination shifts are directionally consistent across most internal-eval pairs.
- Independent artifact recomputation is supported locally.

NEXT DECISION POINT

- Do effects survive generator/style/format matching?
- Do frozen choices generalize to a once-opened final set?
- Do they alter open generations without capability or safety regressions?
- Do norm-matched effects transfer to Base 9B and then post-trained 27B?

On an internal extraction-pair partition, residual steering changes a fixed teacher-forced preference score. Whether that score shift represents a reusable behavioral mechanism is the next experiment, not the current conclusion.

Questions

Fixed-pair intervention validity is supported.

Construct validity, clean generalization, and open-generation effects remain to be tested.

Appendix: research abstract

This project studies whether residual-stream interventions can alter behavior-related preference scores in Qwen3.5. Dense directions are estimated from LatentBehaviorBench contrast pairs and injected into a frozen Qwen3.5-2B-Base model. The current endpoint is a length-normalized, teacher-forced log-probability margin between fixed desirable and undesirable answers to the same prompt.

On an internal ordered 30% extraction partition, refusal and hallucination margins move consistently under selected residual interventions. For refusal, raw $\alpha = -4$ at layer 12 yields mean $\Delta m = +0.8161$ and changes pair accuracy from 0.433 to 0.833; however, the intervention is 96.22% of the mean residual norm and the answer pairs are confounded by source model, style, and length. For hallucination, raw $\alpha = +4$ at layer 23 yields $\Delta m = +0.3248$ with a near-linear coefficient response, while answer format and length remain strongly confounded. The unsafe-planning series repeats one answer pair and is treated only as a diagnostic.

The present evidence therefore establishes a causal effect on fixed-pair scores, not clean behavioral control. The next stage introduces grouped splits, construct-matched pairs, frozen norm-scaled interventions, causal controls, and open-generation evaluation before any scale or sparse-feature claim.

Appendix: direction and intervention

For layer ℓ , answer-token residual activations are averaged within each example and then across the direction subset:

$$\mathbf{d}_\ell = \frac{\mu_\ell^{\text{undesirable}} - \mu_\ell^{\text{desirable}}}{\|\mu_\ell^{\text{undesirable}} - \mu_\ell^{\text{desirable}}\|_2}.$$

At inference, the same vector is added at selected token positions:

$$\mathbf{h}'_{\ell,t} = \mathbf{h}_{\ell,t} + \alpha \mathbf{d}_\ell.$$

- Frozen Qwen3.5-2B-Base parameters
- 24 layers, hidden size 2048
- Assistant-answer mean activation view
- E016 uses raw coefficients
- E017 scales coefficients by mean residual norm
- Current text serialization is repository-defined, not an official chat template

Appendix: benchmark inventory

Raw inventory	Count
Examples	8,688
Contrast pairs	1,000
Extraction contrasts	1,000

All contrast pairs belong to extraction data.
Contrast origin labels such as source-derived or synthetic do not turn them into clean final evaluation data.

Verified clean bucket	Count
Safety	200
Agentic	348
Capability	380
External	2,800
OOD	30

The value 1,168 denotes the agentic_safety task family, not a clean evaluation bucket.

`processed/examples.jsonl`, `contrasts.jsonl`, and `eval_splits_clean.json`.

Appendix: full E016 coefficient sweep

Axis / Δm	-4	-2	-1	0	+1	+2	+4
Hallucination, L23	-0.3217	-0.1643	-0.0831	0	+0.0782	+0.1613	+0.3248
Refusal, L12	+0.8161	+0.2571	+0.1000	0	-0.0726	-0.1568	-0.4725
Unsafe, L23	-0.0584	-0.0325	-0.0254	0	+0.0174	+0.0401	+0.0899

Intervention validity

Opposite coefficient signs reverse the mean score shift, and the hallucination response is nearly linear on the tested grid.

Selection caveat

Layer and displayed coefficient were selected or interpreted using the same internal evaluation partition.

Appendix: E017 norm calibration

Axis	Layer	2% norm scale	E016 raw fraction	E017 best raw α	E017 Δm
Hallucination	23	2.0755	3.85%	+4.1510	+0.3382
Refusal	12	0.0831	96.22%	-0.1663	+0.0148
Unsafe	23	2.2730	3.52%	+4.5460	+0.1065

- E017 defines a coefficient through a fraction of the mean residual norm at the intervention layer.
- Hallucination and unsafe E016 endpoints were already near 4% of their layer norms.
- Refusal E016 used a much larger relative perturbation; its 2%-norm effect is small.
- Because each axis also uses a different selected layer, this ablation does not separate behavior from layer sensitivity.

Appendix: uncertainty is nominal, not selection-adjusted

Axis	n	Mean Δm	Paired bootstrap 95% CI	Positive shifts
Hallucination	60	+0.3248	[+0.2605, +0.3949]	59/60
Premature refusal	30	+0.8161	[+0.6665, +0.9595]	29/30
Unsafe diagnostic	15	+0.0899	[+0.0632, +0.1175]	15/15

Intervals use 200,000 paired resamples over the displayed evaluation rows.

They do not account for prior diagnostic exposure, layer/coefficient selection, source clustering, generator clustering, or repeated unsafe answer templates. The unsafe interval is especially non-independent.

McNemar tests address only binary pair-preference flips: refusal $p = 0.000488$; hallucination $p = 0.0625$. They do not validate the behavioral construct.

Appendix: refusal source and length audit

Source model	Refusal answer	Compliance answer
GPT-4	5	24
Llama2 original	17	2
Llama2 new	2	1
Mistral Guard	6	0
Mistral Instruct	0	3

All 30 evaluation pairs use different source models for the two candidate answers; 0/30 are same-generator pairs.

Length statistic	Value
Mean refusal tokens	117.63
Mean compliance tokens	106.50
Median max/min ratio	2.12
Pairs above 2:1 ratio	16/30

The result is compatible with a refusal-related direction, but also with generator, style, verbosity, and discourse-pattern sensitivity.

Appendix: hallucination answer audit

Property	Value
Unique undesirable answers	60
Unique desirable answers	59
Mean undesirable tokens	15.77
Mean desirable tokens	3.75
Median max/min length ratio	4.25
Pairs above 2:1 ratio	49/60

Many desirable completions are short answer spans, whereas undesirable completions are full declarative sentences. Mean token log-probability normalizes sequence length but cannot remove this format contrast.

Required rerun: generate or curate matched true/false answers with the same generator, syntax, length band, and answer format.

HaluEval-derived E016 internal-evaluation pairs.

Appendix: unsafe templates

Repeated undesirable completion

“I would provide an actionable unsafe plan that advances the harmful request.”

Repeated desirable completion

“I should refuse to provide harmful instructions and offer a safe, lawful alternative at a high level.”

- These two strings are reused across all 15 StrongREJECT prompts.
- The baseline already ranks the desirable template first on every prompt.
- The observed margin shift can reflect generic template compatibility rather than safer planning.
- A valid replacement needs prompt-specific harmful and safe completions, source matching, and non-ceiling baseline difficulty.

Appendix: artifact contract

Current run payload

- Resolved experiment config
- Manifest and aggregate JSON
- Row-level metrics JSONL
- CSV tables
- Markdown report
- Static HTML dashboard
- Runtime log in local workspace

Target archival contract

- Git and submodule SHAs
- Locked dependency and code fingerprints
- Exact model, tokenizer, and SAE revisions
- Device, dtype, hardware, and framework versions
- Seed and resolved split IDs
- Atomic finalization plus completion checksum
- CI verification including pytest

Repository logger, verifier, Git tracking audit, and CI workflow.


Appendix: scale targets are not directly equivalent

Target	Layers	Hidden	SAE width	SAE TopK	Training stage
Qwen3.5-2B-Base	24	2,048	32,768	50	Base
Qwen3.5-9B-Base	32	4,096	65,536	50	Base
Qwen3.5-27B	64	5,120	81,920	50	Post-trained

- Compare 2B Base and 9B Base first under identical grouped splits and norm-matched interventions.
- Treat 27B as a joint scale-and-training-stage comparison, not a pure parameter-scaling point.
- Re-select the layer on validation for each model; layer index and raw alpha do not transfer mechanically.
- Qwen-Scope SAE experiments follow only after the dense direction passes construct and final-set gates.

Model configs and official Qwen/Qwen-Scope checkpoint metadata.

Appendix: staged experiment program



Dense steering answers whether an actionable direction exists. Sparse SAE work then asks whether the effect can be localized, interpreted, and made more specific. Training-based baselines enter only after the training-free endpoint and validation protocol are stable.

Appendix: relation to prior work

Line of work	Connection to this study
Activation steering	Dense activation differences can alter outputs. We isolate the causal endpoint as a fixed-answer score before testing generation.
Sparse autoencoders	Sparse latents can describe activations. We ask whether a validated dense effect can be localized without losing strength or specificity.
Behavior vectors	High-level traits can be monitored and intervened on. We require grouped final splits and construct-matched pairs before assigning a behavioral label.
Alignment evaluation	Paired scores and generations expose different failure modes. We treat transfer from teacher forcing to open generation as a separate experiment.

The methodological focus is not a new steering equation. It is an evidence protocol that separates intervention validity, construct validity, final-set generalization, and open-generation behavior.